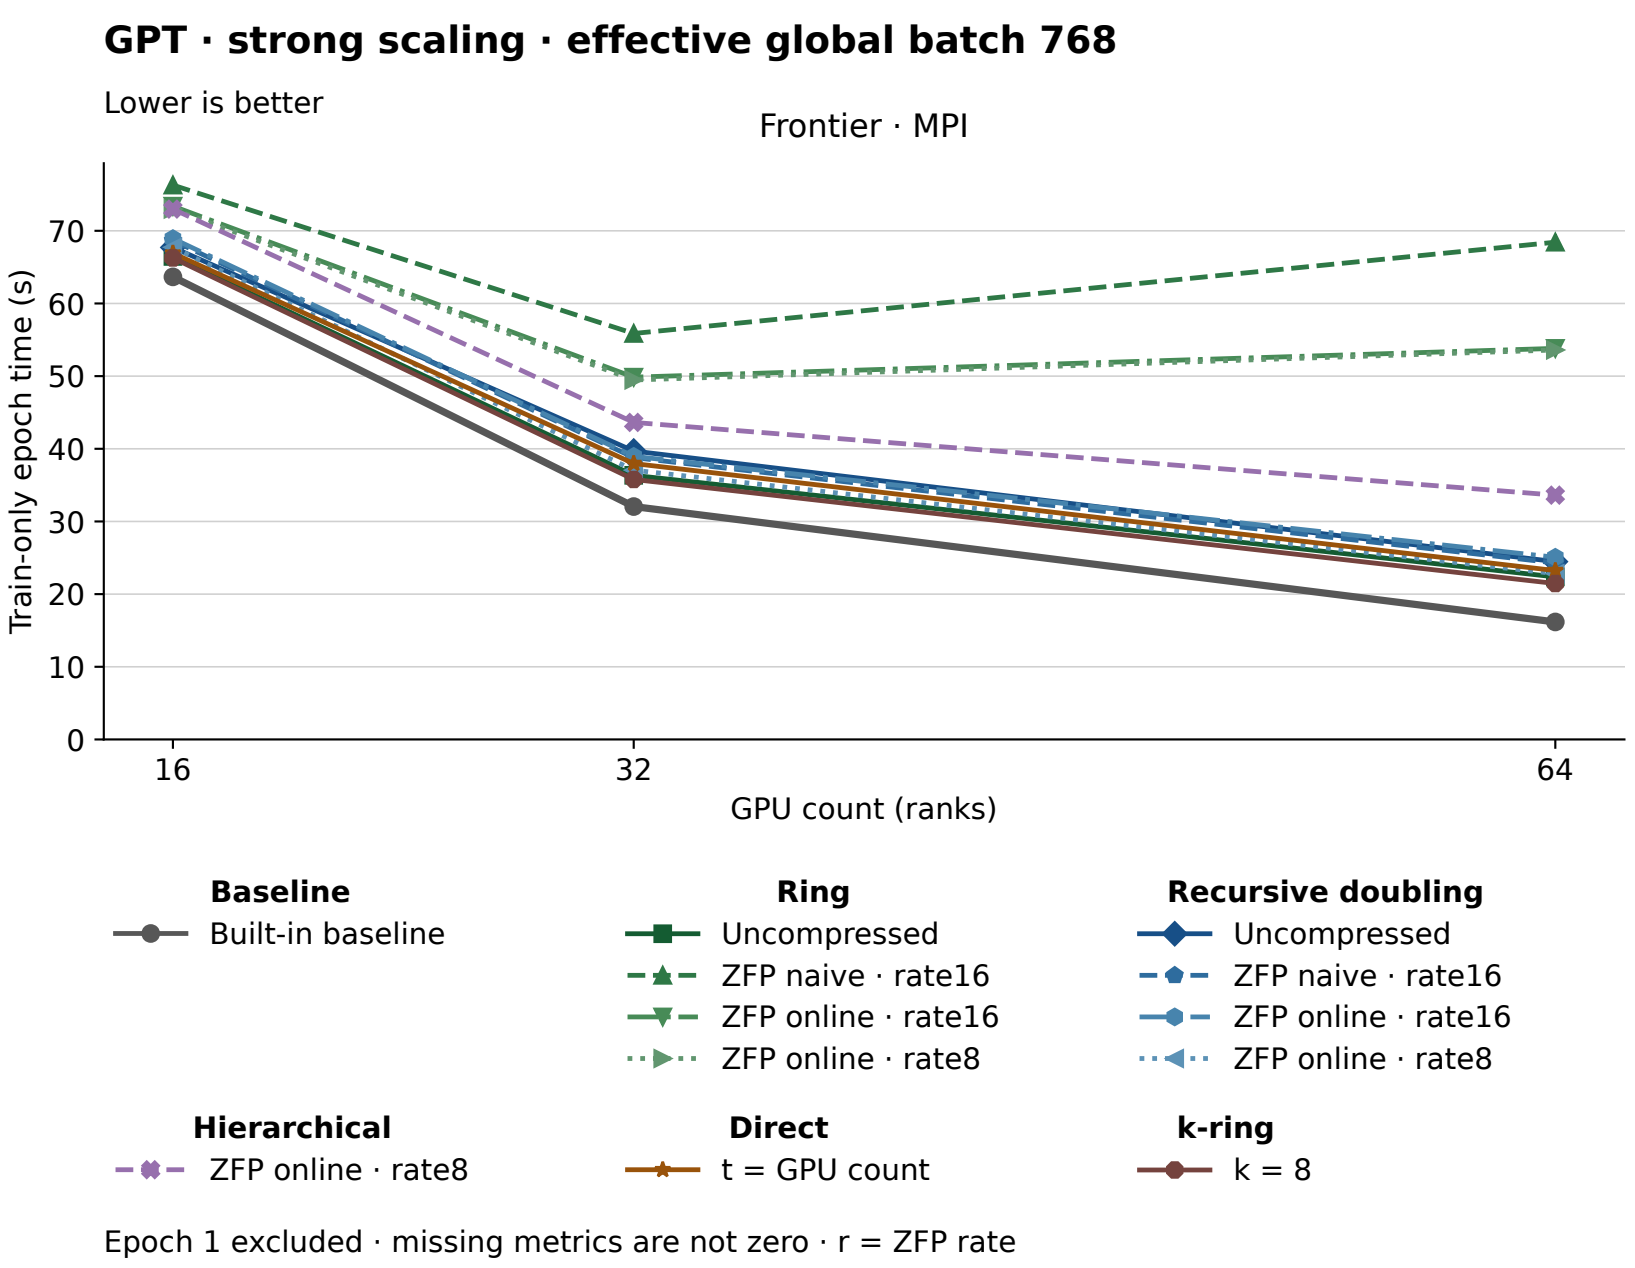
GPT · strong scaling · effective global batch 768
Lower is better
Frontier · MPI
Train-only epoch time (s)
0
10
20
30
40
50
60
70
16
32
64
GPU count (ranks)
Baseline
Built-in baseline
Ring
Uncompressed
ZFP naive · rate16
ZFP online · rate16
ZFP online · rate8
Recursive doubling
Uncompressed
ZFP naive · rate16
ZFP online · rate16
ZFP online · rate8
Hierarchical
ZFP online · rate8
Direct
t = GPU count
k-ring
k = 8
Epoch 1 excluded · missing metrics are not zero · r = ZFP rate